



PennState

CMPSC 311 - Introduction to Systems Programming

Bit/Byte Operations

Professor Abutalib Aghayev

(Slides are mostly by Professor Patrick McDaniel)



Hobbit

Hobbyte

Number Systems



- All base-X systems have the following characteristic:

Assume a base b and digits $P = \{p_k, p_{k-1}, p_{k-2}, \dots, p_1, p_0\}$

$$value = \sum_{i=0}^k b^i * p_i$$

where $\forall p_i \in P, p_i = [0, b - 1]$

Base-10 (decimal)

- Digits: 0,1,2,3,4,5,6,7,8,9
- Place values: $10^0, 10^1, 10^2, \dots$
- Example: $123 = 3 \times 1 + 2 \times 10 + 1 \times 100$
 $= 123$

Base-2(binary)

- Digits: 0,1
- Place values: $2^0, 2^1, 2^2, \dots$
- Example: $1011 = 1 \times 1 + 1 \times 2 + 0 \times 4 + 1 \times 8$
 $= 11$

Base-16 (hexadecimal)

- Digits: 0,1,2,3,4,5,6,7,8,9,A,B,C,D,E,F
- Place values: $16^0, 16^1, 16^2, \dots$
- Example: $0xAFC = C \times 1 + F \times 16 + A \times 256$
 $= 2812$

Converting Decimal to Binary

- while $n \neq 0$:
 - next binary digit (from right to left) = $n \% 2$
 - $n = n / 2$

n	remainder	digit #
235	1	0
117	1	1
58	0	2
29	1	3
14	0	4
7	1	5
3	1	6
1	1	7

235 (base-10) = 11101011 (base-2)

Converting Decimal and Binary to Hex



- Converting decimal to hexadecimal
 - while $n \neq 0$:
 - next hex digit (from right to left) = $n \% 16$
 - $n = n / 16$

n	remainder	digit #
235	B (11)	0
14	E (14)	1

235 (base-10) = EB (base-16)

- Converting binary to hexadecimal
 - convert decimal to binary
 - group binary digits to into groups of 4 bits (nibbles) starting from right
 - Convert each nibble to hexadecimal digit

235 (base-10) = 11101011 (base-2) = 1110 1011 = EB

- Be familiar with hex notation: 0xDEADBEEF has 8 nibbles, 4 bytes (or octets)
- Hexadecimal to binary: just convert each nibble to binary

Conversion Summary




- Binary to decimal
 - sum of binary digits times powers of 2
- Hexadecimal to decimal
 - sum of hex digits times power of 16
- Decimal to binary
 - division method
- Decimal to hex
 - division method
 - or, convert to binary and use binary to hex method below
- Binary to hex
 - group binary digits to nibbles, convert nibbles to hex digits
- Hex to binary
 - convert each hex digit to binary



Byte Ordering Example

- How should bytes within a multi-byte word be ordered in memory?
 - Big Endian: Least significant byte has *highest* address (SPARC)
 - Little Endian: Least significant byte has *lowest* address (x86)
- Example: `int x = 305419896`
 - Variable `x` has 4-byte representation `0x12345678`
 - Address given by `&x` is `0x100`

Big Endian



Little Endian



Examining Data Representations



- Code to find endianness of the architecture
 - Casting pointer to `uint8_t *` creates byte array

```
#include <stdio.h>
#include <stdint.h>

void show_bytes(uint8_t *start, int len) {
    for (int i = 0; i < len; ++i)
        printf("%p\t0x%.2x\n", start+i, start[i]);
    printf("\n");
}

int main(void) {
    int a = 305419896;
    printf("a lives at address %p\n\n", &a);
    show_bytes((uint8_t *)&a, sizeof(int));
}
```

Result (Linux on x86):

a lives at address 0x7ffebf803174

0x7ffebf803174	0x78
0x7ffebf803175	0x56
0x7ffebf803176	0x34
0x7ffebf803177	0x12

printf directives

%p Print pointer

%x Print hexadecimal

Boolean Algebra

- Developed by George Boole in 19th Century
 - Algebraic representation of logic that based on “True” (as 1) and “False” (as 0)

And

$A \& B = 1$ when both $A=1$ and $B=1$

$\&$	0	1
0	0	0
1	0	1

Or

$A | B = 1$ when either $A=1$ or $B=1$

$ $	0	1
0	0	1
1	1	1

Not

$\sim A = 1$ when $A=0$

\sim	
0	1
1	0


Exclusive-Or (Xor)

$A \wedge B = 1$ when either $A=1$ or $B=1$, but not both

\wedge	0	1
0	0	1
1	1	0

Bit-Level Operations in C

- Operations `&`, `|`, `~`, `^` Available in C
 - Apply to any “integral” data type
 - long, int, short, char, unsigned
 - View arguments as bit vectors
 - Arguments applied bit-wise
- Examples (Char data type)
 - `~0x41` \rightarrow `0xBE`
 - `~010000012` \rightarrow `101111102`
 - `~0x00` \rightarrow `0xFF`
 - `~000000002` \rightarrow `111111112`
 - `0x69 & 0x55` \rightarrow `0x41`
 - `011010012 & 010101012` \rightarrow `010000012`
 - `0x69 | 0x55` \rightarrow `0x7D`
 - `011010012 | 010101012` \rightarrow `011111012`



Contrast: Logic Operations in C

- Contrast to Logical Operators (`&&`, `||`, `!`)
 - View 0 as “False”
 - Anything nonzero as “True”
 - Always return 0 or 1
 - Early termination

Representing & Manipulating Sets PennState

- Representation

- Width w bit vector represents subsets of $\{0, \dots, w-1\}$ for set “A”
- $a_j = 1$ if $j \in A$

01010101 { 0, 2, 4, 6 }
76543210

01101001 { 0, 3, 5, 6 }
76543210

Operations On Sets:

& Intersection	01000001	{ 0, 6 }
Union	01111101	{ 0, 2, 3, 4, 5, 6 }
^ Symmetric difference	00111100	{ 2, 3, 4, 5 }
~ Complement	10101010	{ 1, 3, 5, 7 }

Representing Signed Numbers

b3	b2	b1	b0		Unsigned	One's complement	Two's complement
0	0	0	0		0	0	0
0	0	0	1		1	1	1
0	0	1	0		2	2	2
0	0	1	1		3	3	3
0	1	0	0		4	4	4
0	1	0	1		5	5	5
0	1	1	0		6	6	6
0	1	1	1		7	7	$7 = 2^{n-1} - 1$
1	0	0	0		8	-7	$-8 = -2^{n-1}$
1	0	0	1		9	-6	-7
1	0	1	0		10	-5	-6
1	0	1	1		11	-4	-5
1	1	0	0		12	-3	-4
1	1	0	1		13	-2	-3
1	1	1	0		14	-1	-2
1	1	1	1		15	-0	-1

- Computing one's complement negative representation:
 - Complement the positive number
- Computing two's complement negative representation:
 - Complement the positive number and add 1
- Most architectures use two's complement
- Given n-bit signed integer:
 - Positive range: $0 - 2^{n-1} - 1$
 - Negative range: -2^{n-1}

type	bytes (32-bit)	bytes (64-bit)	32-bit range	printf
char	1	1	[0, 255]	%c
short int	2	2	[-32768, 32767]	%hd
unsigned short int	2	2	[0, 65535]	%hu
int	4	4	[-2147483648, 2147483647]	%d
unsigned int	4	4	[0, 4294967295]	%u
long int	4	8	[-2147483648,	%ld

Beware of integer overflows

- Spot the bug in the following:

```
for (uint32_t n = 10; n >=0; --n) {  
    printf("do I even terminate?\n");  
}
```

- What does the following print?

```
int x = 0xffffffff;  
printf("%d\n", x);
```



The latest news from Google AI

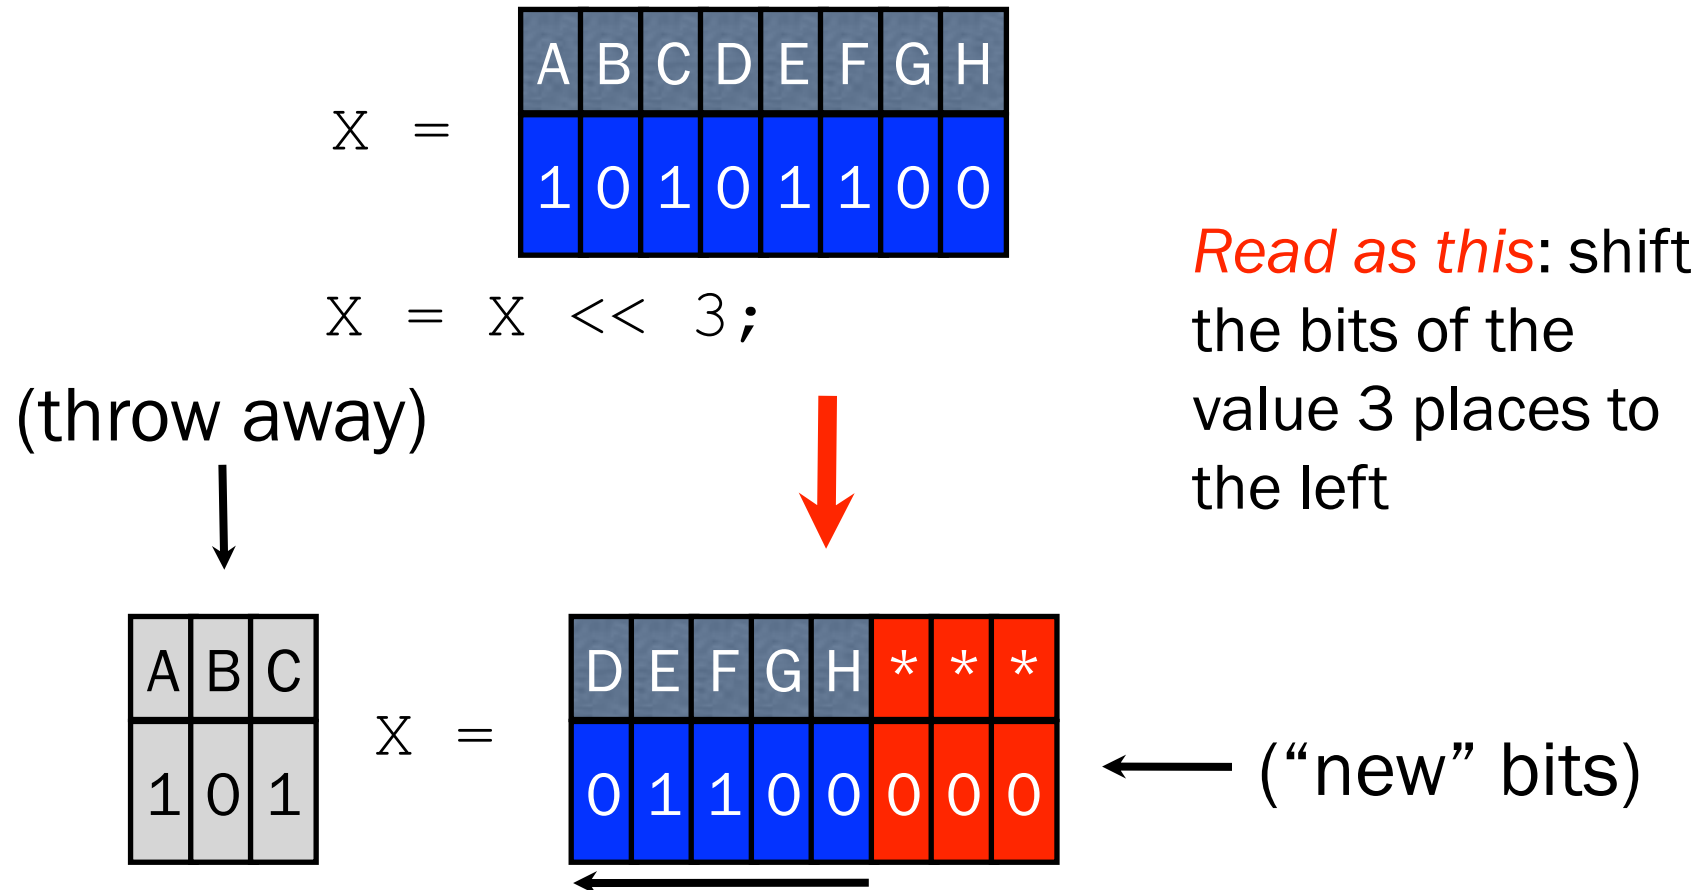
Extra, Extra - Read All About It: Nearly All Binary Searches and Mergesorts are Broken

Friday, June 2, 2006

Posted by Joshua Bloch, Software Engineer

Shift Operations

- A shift operator (\ll or \gg) moves bits to the right or left, throwing away bits and adding bits as necessary



Putting it all together

- Suppose you want to place multiple values in the same 32-bit integer
 - Value **a** in least significant byte
 - Value **b** in 2nd byte
 - Value **c** in 3rd byte
 - Value **d** in 4th byte

Bits	31-24	23-16	8-15	0-7
Values	d	c	b	a

Using bit operations ...



```
uint32_t pack_bytes(uint32_t a, uint32_t b, uint32_t c, uint32_t d) {  
  
    // Setup some local values  
    uint32_t retval = 0x0, tempa, tempb, tempc, tempd;  
  
    tempa = a&0xff; // Make sure you are only getting the bottom 8 bits  
    tempb = (b&0xff) << 8; // Shift value to the second byte  
    tempc = (c&0xff) << 16; // Shift value to the third byte  
    tempd = (d&0xff) << 24; // Shift value to the top byte  
    retval = tempa|tempb|tempc|tempd; // Now combine all of the values  
  
    // Print out all of the values  
    printf("A: 0x%08x\n", tempa);  
    printf("B: 0x%08x\n", tempb);  
    printf("C: 0x%08x\n", tempc);  
    printf("D: 0x%08x\n", tempd);  
  
    // Return the computed value  
    return retval;  
}  
  
...  
  
printf("Packed bytes : 0x%08x\n", pack_bytes(0x111, 0x222, 0x333, 0x444));
```

A: 0x00000011
B: 0x00002200
C: 0x00330000
D: 0x44000000
Packed bytes : 0x44332211